

# Market Sentiment Classifier

Group: Bimal Rajbhandari, Kathleen Nie, Nick Pandolfi

ADSP 32021 ON01 Machine Learning Operations

Term: Fall 2025



# Agenda

1. Business Problem
2. Data Management & Preparation
3. Exploratory Data Analysis
4. Pipeline & Modeling
5. Deployment
6. Monitoring
7. Summary



# Predicting Market Sentiment

## Why It Matters

- Market sentiment (fear/neutral/greed) drives short-term trading and risk exposure.
- Early prediction enables better portfolio adjustments and risk management.

## How the Model Works

- **Inputs:**
  - SPY/VIX: Price, volatility, moving averages
  - News headlines: Sentiment scores (VADER, TextBlob)
- **Output:** Predicts tomorrow's sentiment (fear/neutral/greed)

# MLOPS Component and Lifecycle





# Data Management & Preparation

# Data Preparation

## Data Source

- Fetched historical SPY (S&P 500 proxy) and VIX (Volatility Index) data from [Yahoo Finance API \(yfinance\)](#).
- Financial news headlines from Kaggle
- *Time Horizon:* 2008-01-02 to 2024-03-04.

## Feature Engineering & Target Creation

We transformed raw inputs into 11 predictive features and the target variable

**Time-Series Lagging:** All predictor variables were explicitly lagged by one day to ensure the model uses only t-1 information for predicting.

**Target Labeling:** The dependent variable, [Sentiment\\_Label](#), was engineered by applying rules-based thresholds to the VIX Close price (VIX > 30 = Fear, VIX < 20 = Greed, else Neutral).

# Data Preparation

Processed and aggregated news headline sentiment

- Applied VADER compound scores
- Applied TextBlob polarity and subjectivity

Engineered 11 Lagged Features (using only t-1 data)

– VIX-derived features:

- Lag\_1\_VIX\_Close
- Lag\_1\_VIX\_Daily\_Change
- Lag\_1\_VIX\_7D\_MA
- Lag\_1\_VIX\_30D\_MA

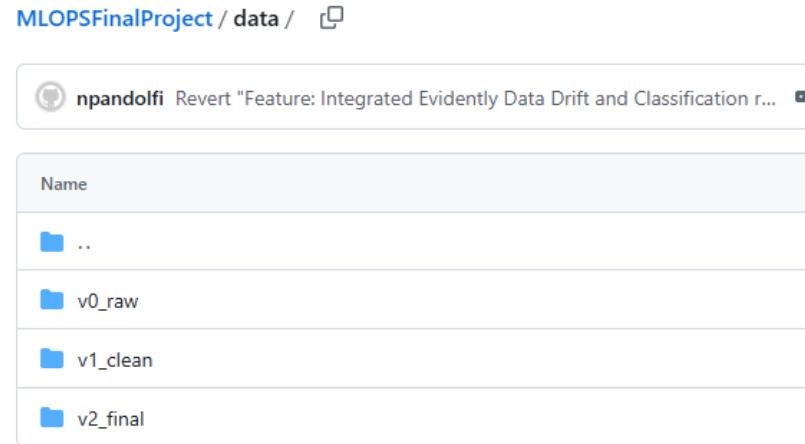
– SPY-derived features:

- Lag\_1\_SPY\_Close
- Lag\_1\_SPY\_Daily\_Return
- Lag\_1\_SPY\_7D\_Return
- Lag\_1\_SPY\_30D\_Return

# Data Management

## Data Versioning & Structure

- Raw market and news un-processed data stored as Version 0
- Merged Feature engineered SPY, VIX and news dataset as Version 1
- Feature-engineered news,SPY, VIX modeling-ready dataset as Version 2
- **Git Integration:** All versions tracked using GitHub for reproducibility and data management





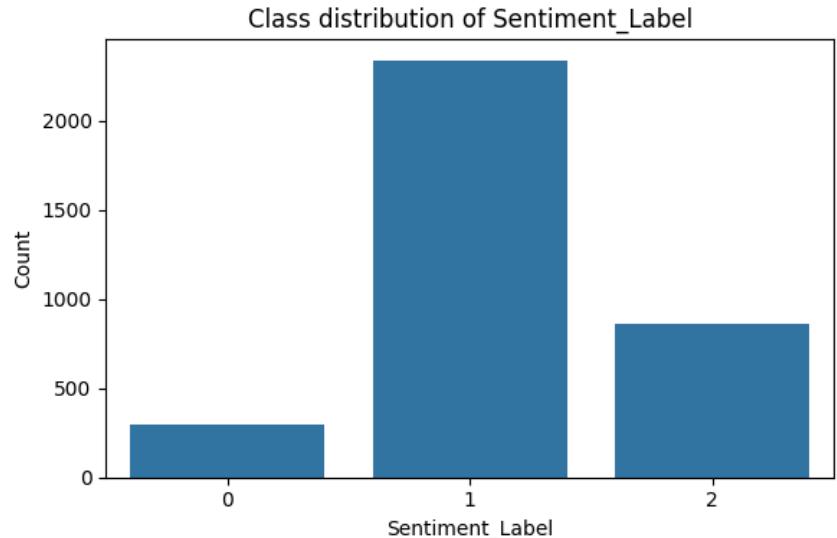
# Exploratory Data Analysis

# Exploratory Data Analysis

We use a mix of Quantitative & Qualitative Data to build our model

- **Yahoo Finance:** API that allows deep, long-term access to Stock Market Data for quantitative purposes. We pull VIX and S&P500 daily returns
- **Financial News:** S&P 500 with Financial News Headlines (2008–2024) dataset from Kaggle to extract qualitative data

These datasets are combined in the Data Engineering Phase to build a final Dataset

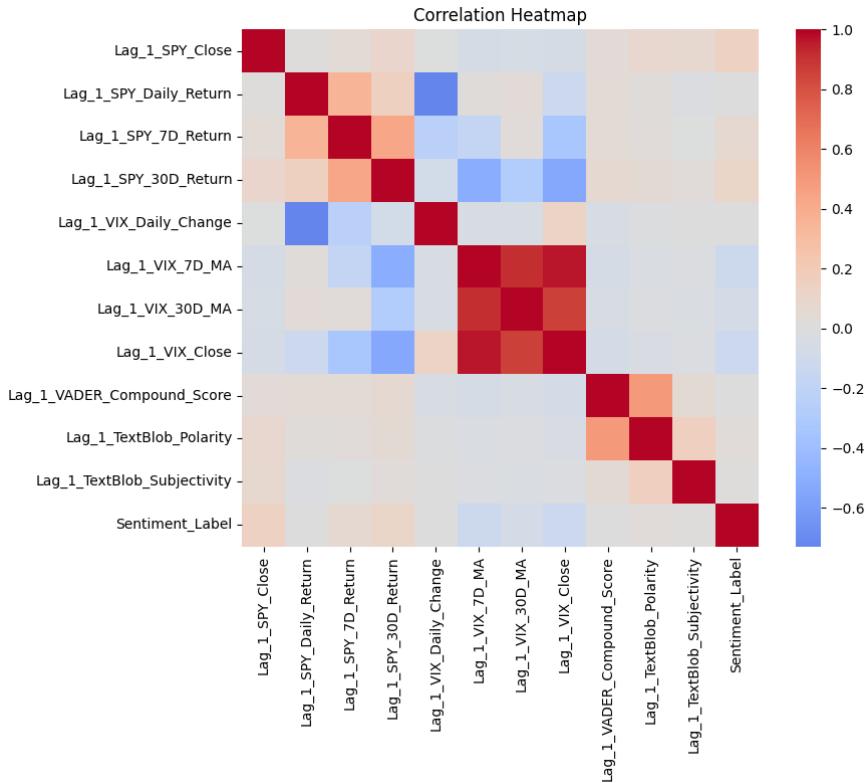


# Exploratory Data Analysis

## Correlation Structure of Market & Sentiment Features

- VIX group is strongly intercorrelated
- SPY-related features form their own block
- Text sentiment is weakly correlated but non-zero

Sentiment\_Label shows meaningful relationships to SPY/VIX changes but weaker ties to text scores.



# Exploratory Data Analysis

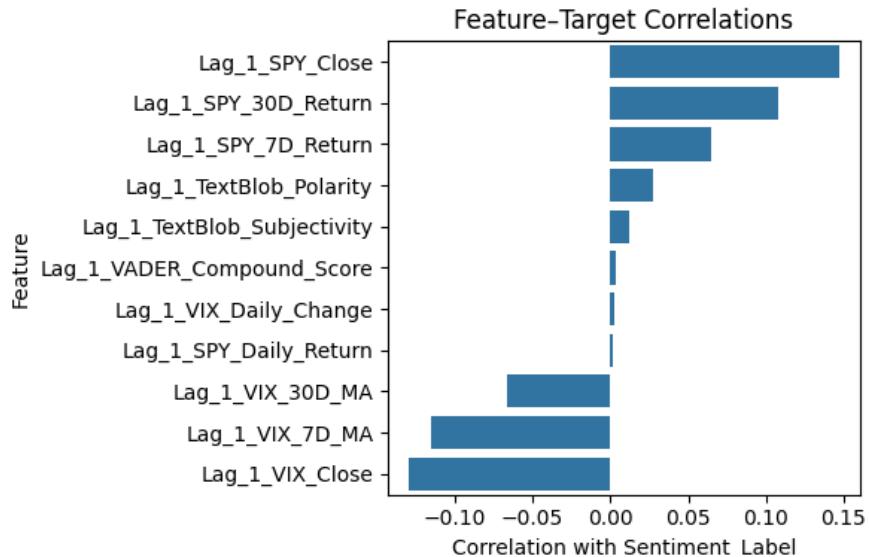
## Key Predictive Features for Market Sentiment

Top positive-leverage features:

- Lag\_1\_SPY\_Close
- Lag\_1\_SPY\_7D\_Return
- Lag\_1\_SPY\_30D\_Return

Top negative-leverage features:

- Lag\_1\_VIX\_Close
- Lag\_1\_VIX\_7D\_MA
- Lag\_1\_VIX\_30D\_MA





# Pipeline & Modeling

# Scikit-Learn Pipeline

## Handling Class Imbalance

### Pipeline 1 – FLAML with Class Weighting

- Used FLAML's `AutoML` with `class_weight='balanced'` (via custom function).
- Class weighting forces the model to pay more attention to minority classes.
- Models: Random Forest, XGBoost, LightGBM (tree-based, support class weights).
- Metric: F1-score (better for imbalanced data than accuracy).

### Pipeline 2 – FLAML with SMOTE

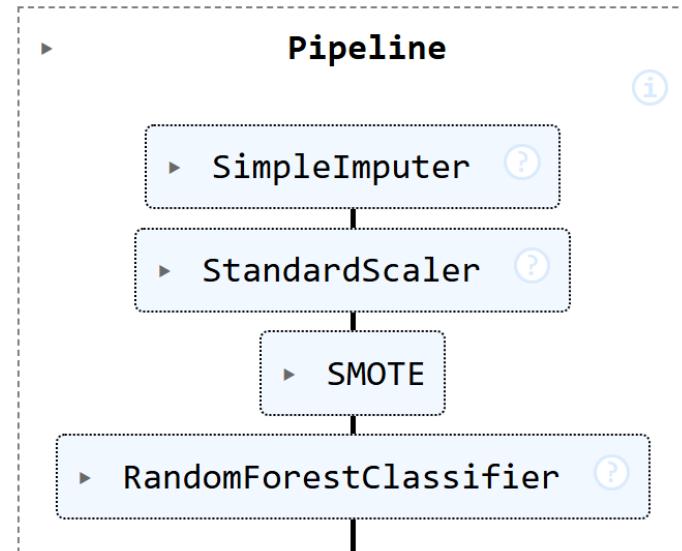
- Used `ImbPipeline` to integrate SMOTE into the workflow.
- Steps: Standard scaling → SMOTE → FLAML AutoML.
- SMOTE creates synthetic samples, balancing class distribution.

# Scikit-Learn Pipeline

## Handling Class Imbalance

### Pipeline 3 - FLAML Fear-Weighted Random Forest with SMOTE

- Built on FLAML AutoML's best RandomForest hyperparameters
- Incorporated Fear-focused class weights
- Combination of pipeline 1 and 2



# Scikit-Learn Pipeline

Estimator	Log Loss	Training Time (s)
Random Forest	0.2555	47.6
LightGBM	0.2075	50.0
XGBoost	0.2069	139.0
Extra Trees	<b>0.1947 (Best)</b>	187.9
XGB (Limited Depth)	0.2149	190.4
SGD	0.3925	190.9
L1 Logistic	0.1978	191.6

## Pipeline 1 – FLAML with Class Weighting: Training Summary

- **Best Model: Extra Trees Classifier**  
(log loss = 0.1947).
- **Final Model Parameters:**
  - `criterion='entropy'`
  - `max_features=1.0`
  - `max_leaf_nodes=18344`
  - `n_estimators=2047`
  - `n_jobs=-1`

# Scikit-Learn Pipeline

Estimator	Log Loss	Training Time (s)
Random Forest	0.1571	104.4
LightGBM	0.1477	108.0
XGBoost	0.1634	206.9
Extra Trees	<b>0.1104 (Best)</b>	278.8
XGB (Limited Depth)	0.1622	282.8
SGD	0.5306	283.7
L1 Logistic	0.2105	285.0

## Pipeline 2 – FLAML with SMOTE

- **Best Model: Extra Trees Classifier** (log loss = 0.1104).
- **Final Model Parameters:**
  - `criterion='entropy'`
  - `max_features=1.0`
  - `max_leaf_nodes=18344`
  - `n_estimators=2047`
  - `n_jobs=-1`

# Scikit-Learn Pipeline

## Maximizing Fear Detection for Portfolio Safety

### SMOTE Model – Better for Risk Management

- **Fear (Class 0)** = Critical, rare signal to "go defensive" and avoid downside risk.
- **False Negatives (missing Fear)** = Staying aggressively invested during market downturns → **high risk**.
- **SMOTE captures 3% more Fear instances** (83% vs. 80% recall).
- **Prioritize safety**: Higher recall reduces costly missed signals, even with a slight precision trade-off



Metric	SMOTE Model	Class Weight Model
<b>Fear Recall</b>	83%	80%
<b>True Positives</b>	50	48

# Deployment



# Dockerization for Reproducible Model Deployment

**Objective:** To package the trained model and its dependencies into a single, portable, and reproducible unit for easy deployment.

## Dockerization Process:

1. **Dependencies:** A `requirements.txt` file listed all necessary Python libraries (FLAML, Scikit-learn, FastAPI, uvicorn, etc.).
2. **Docker Structure:** The Dockerfile contained instructions to:
  - Set up a base Python environment.
  - Install the required dependencies.
  - Copy the saved model pipeline file.
  - Copy the FastAPI application code (`api.py`).
3. **FastAPI Endpoint** (lightweight, high-performance).
  - Accepts market and sentiment features as input.
  - Returns predicted sentiment class (`0`, `1`, or `2`).

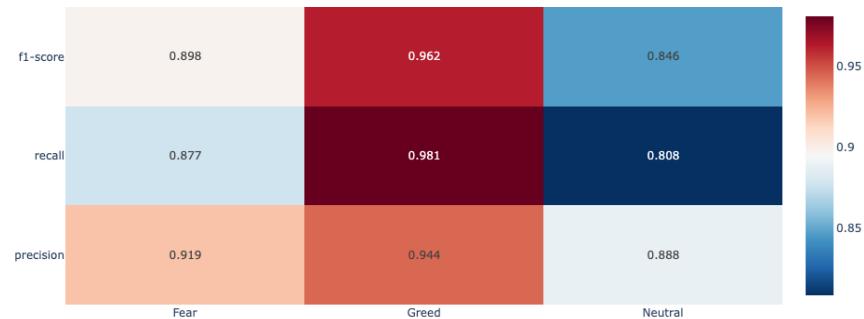
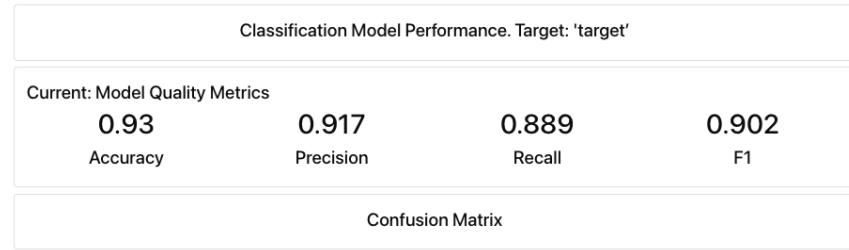
# Model Monitoring With Evidently



# Model Monitoring with Evidently

## Classification Evaluation:

- **Evidently Performance Tracking** shows the accuracy of our classification model at 93% accuracy
- **Multiclass Classification:** Our model performs extremely well at predicting **Greed** and the most poorly at predicting **Neutral**
- Evidently also provides a data summary by evaluating distributions, counts, NAs, etc

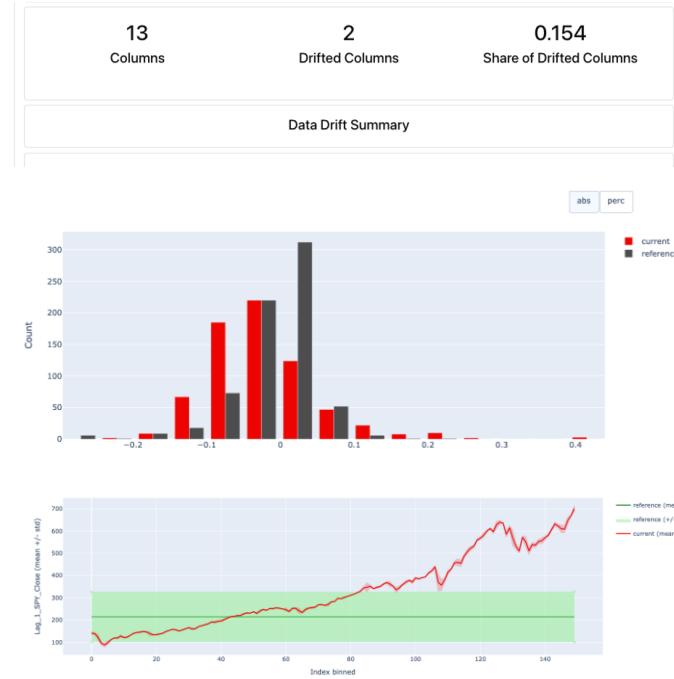


# Data Drift

## Real-World Data Drift Scenarios:

- **Regime Change:** We simulate a regime change by multiplying the daily S&P500 by 120% (organic stock market rise of 20% in 2 years). This highlights the continuous need to re-train our model as time goes on
- **Feature Swap:** We also simulate an error of 2 columns being mixed up in the data ingestion process (Lag\_1\_SPY\_30D\_Return and Lag\_1\_VIX\_Daily\_Change)

Our Model Monitoring is able to detect both of these scenarios via the K-S p\_value (test for difference in distributions), highlighting 2 drifted columns and 15% share of drifted columns



# Summary

- **Recap**

Models are often developed and proven effective in a sterile research environment (the laptop), but they fail to be deployed, scaled, or reliably maintained in the real-world production environment without MLOps

- **Pipeline & Engineering**

Our data versioning provides reproducible snapshots of data and features used for training, and pipelines automate the end-to-end workflow using our versioned assets.

- **Model Overview**

AutoML can help identify the best performing model for our given dataset to redistribute human resources to other areas

- **Model Monitoring**

Evidently can be great for automated performance tracking and data drift monitoring to ensure our model is consistently being refreshed as data shifts for various reasons

**Thank you!**



# References

---

Check it out on GitHub!

<https://github.com/rajbhanb/MLOPSFinalProject>

<https://www.kaggle.com/datasets/notlucasp/financial-news-headlines>